

Problem Statement

Linear regression by using Deep Neural network: Implement Boston housing price prediction problem by Linear regression using Deep Neural network. Use Boston House price prediction dataset.

Import Library

```
In [65]: # Data analysis and visualization
import tensorflow as tf
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline

# Preprocessing and evaluation
from sklearn.model_selection import train_test_split
from sklearn.compose import make_column_transformer
from sklearn.preprocessing import MinMaxScaler
```

Load Data

```
In [66]: (X_train , y_train), (X_test , y_test) = tf.keras.datasets.boston_housing.load_data(
                                                path = 'boston_housing_npz',
                                                test_split = 0.2,
                                                seed = 42
                                                )
```

```
In [67]: # Checking the data shape and type
(X_train.shape, type(X_train)), (X_test.shape, type(X_test)), (y_train.shape, ty
```

```
Out[67]: (((404, 13), numpy.ndarray),
          ((102, 13), numpy.ndarray),
          ((404,), numpy.ndarray),
          ((102,), numpy.ndarray))
```

```
In [68]: # Converting Data to DataFrame
X_train_df = pd.DataFrame(X_train)
y_train_df = pd.DataFrame(y_train)

# Preview the training data
X_train_df.head(10)
```

Out[68]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.09178	0.0	4.05	0.0	0.510	6.416	84.1	2.6463	5.0	296.0	16.6	395.50	9.04
1	0.05644	40.0	6.41	1.0	0.447	6.758	32.9	4.0776	4.0	254.0	17.6	396.90	3.53
2	0.10574	0.0	27.74	0.0	0.609	5.983	98.8	1.8681	4.0	711.0	20.1	390.11	18.07
3	0.09164	0.0	10.81	0.0	0.413	6.065	7.8	5.2873	4.0	305.0	19.2	390.91	5.52
4	5.09017	0.0	18.10	0.0	0.713	6.297	91.8	2.3682	24.0	666.0	20.2	385.09	17.27
5	0.10153	0.0	12.83	0.0	0.437	6.279	74.5	4.0522	5.0	398.0	18.7	373.66	11.97
6	0.31827	0.0	9.90	0.0	0.544	5.914	83.2	3.9986	4.0	304.0	18.4	390.70	18.33
7	0.29090	0.0	21.89	0.0	0.624	6.174	93.6	1.6119	4.0	437.0	21.2	388.08	24.16
8	4.03841	0.0	18.10	0.0	0.532	6.229	90.7	3.0993	24.0	666.0	20.2	395.33	12.87
9	0.22438	0.0	9.69	0.0	0.585	6.027	79.7	2.4982	6.0	391.0	19.2	396.90	14.33

In [69]: *# View summary of datasets*
`X_train_df.info()`
`print('_', '*40)`
`y_train_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 404 entries, 0 to 403
Data columns (total 13 columns):
#   Column  Non-Null Count  Dtype
---  -
0    0      404 non-null    float64
1    1      404 non-null    float64
2    2      404 non-null    float64
3    3      404 non-null    float64
4    4      404 non-null    float64
5    5      404 non-null    float64
6    6      404 non-null    float64
7    7      404 non-null    float64
8    8      404 non-null    float64
9    9      404 non-null    float64
10   10     404 non-null    float64
11   11     404 non-null    float64
12   12     404 non-null    float64
dtypes: float64(13)
memory usage: 41.2 KB
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 404 entries, 0 to 403
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0    0      404 non-null    float64
dtypes: float64(1)
memory usage: 3.3 KB
```

In [70]: `X_train_df.describe()`

Out[70]:

	0	1	2	3	4	5	6	
count	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404
mean	3.789989	11.568069	11.214059	0.069307	0.554524	6.284824	69.119307	3
std	9.132761	24.269648	6.925462	0.254290	0.116408	0.723759	28.034606	2
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1
25%	0.081960	0.000000	5.190000	0.000000	0.452000	5.878750	45.475000	2
50%	0.262660	0.000000	9.690000	0.000000	0.538000	6.210000	77.500000	3
75%	3.717875	12.500000	18.100000	0.000000	0.624000	6.620500	94.425000	5
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12

Preprocessing

```

In [71]: # Create column transformer
ct = make_column_transformer(
    (MinMaxScaler(), [0, 1, 2, 4, 5, 6, 7, 8, 9, 10, 11, 12])
)

# Normalization and data type change
X_train = ct.fit_transform(X_train).astype('float32')
X_test = ct.transform(X_test).astype('float32')
y_train = y_train.astype('float32')
y_test = y_test.astype('float32')

# Distribution of X_train feature values after normalization
pd.DataFrame(X_train).describe()

```

Out[71]:

	0	1	2	3	4	5	6	
count	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404.000000	404
mean	0.042528	0.115681	0.394210	0.348815	0.521905	0.681970	0.241618	0
std	0.102650	0.242696	0.253866	0.239522	0.138678	0.288719	0.194973	0
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0
25%	0.000850	0.000000	0.173387	0.137860	0.444098	0.438466	0.087361	0
50%	0.002881	0.000000	0.338343	0.314815	0.507569	0.768280	0.184767	0
75%	0.041717	0.125000	0.646628	0.491770	0.586223	0.942585	0.362255	1
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1

Model, Predict, Evaluation

```

In [72]: X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.
X_train.shape, X_val.shape, y_train.shape, y_val.shape

```

Out[72]: ((363, 12), (41, 12), (363,), (41,))

Creating the Model and Optimizing the Learning Rate learning rate = 0.01, batch_size = 32, dense_layers = 2, hidden_units for Dense_1 layer= 10, hidden_units for Dense_2 layer = 100


```
In [73]: # Set random seed
tf.random.set_seed(42)


# Building the model
model = tf.keras.Sequential([
    tf.keras.layers.Dense(units=10, activation='relu', input_shape=(X_train.shape[
    tf.keras.layers.Dense(units=100, activation='relu', name='Dense_2'),
    tf.keras.layers.Dense(units=1, name='Prediction')
])


# Compiling the model
model.compile(
    loss = tf.keras.losses.mean_squared_error,
    optimizer = tf.keras.optimizers.RMSprop(learning_rate=0.01),
    metrics = ['mse']
)


# Training the model
history = model.fit(
    X_train,
    y_train,
    batch_size=32,
    epochs=50,
    validation_data=(X_val, y_val)
)
```


```
C:\Users\admin\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras
\src\layers\core\dense.py:88: UserWarning: Do not pass an `input_shape`/`input_
dim` argument to a layer. When using Sequential models, prefer using an `Input
(shape)` object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```


Epoch 1/50
12/12  2s 32ms/step - loss: 373.0555 - mse: 373.3045 - val_loss: 95.6019 - val_mse: 125.1949


Epoch 2/50
12/12  0s 8ms/step - loss: 90.9924 - mse: 90.7250 - val_loss: 59.1337 - val_mse: 75.6957


Epoch 3/50
12/12  0s 9ms/step - loss: 57.1830 - mse: 57.1047 - val_loss: 48.6360 - val_mse: 58.5284


Epoch 4/50
12/12  0s 7ms/step - loss: 47.4303 - mse: 47.4261 - val_loss: 45.1271 - val_mse: 52.3747


Epoch 5/50
12/12  0s 8ms/step - loss: 43.4081 - mse: 43.4368 - val_loss: 41.5549 - val_mse: 47.4119


Epoch 6/50
12/12  0s 7ms/step - loss: 39.8037 - mse: 39.8562 - val_loss: 36.3672 - val_mse: 41.5954


Epoch 7/50
12/12  0s 8ms/step - loss: 35.7645 - mse: 35.8292 - val_loss: 31.5579 - val_mse: 36.1743


Epoch 8/50
12/12  0s 9ms/step - loss: 32.1831 - mse: 32.2583 - val_loss: 27.7860 - val_mse: 31.6564


Epoch 9/50
12/12  0s 8ms/step - loss: 29.3124 - mse: 29.3902 - val_loss: 25.0090 - val_mse: 28.5063


Epoch 10/50
12/12  0s 9ms/step - loss: 27.2456 - mse: 27.3230 - val_loss: 22.5495 - val_mse: 25.6590


Epoch 11/50
12/12  0s 9ms/step - loss: 25.5360 - mse: 25.6209 - val_loss: 20.9782 - val_mse: 23.9131


Epoch 12/50
12/12  0s 12ms/step - loss: 24.5268 - mse: 24.6152 - val_loss: 19.6550 - val_mse: 22.4570


Epoch 13/50
12/12  0s 10ms/step - loss: 23.6689 - mse: 23.7587 - val_loss: 18.7873 - val_mse: 21.3226


Epoch 14/50
12/12  0s 8ms/step - loss: 22.9617 - mse: 23.0537 - val_loss: 17.9019 - val_mse: 20.1441


Epoch 15/50
12/12  0s 7ms/step - loss: 22.4037 - mse: 22.4961 - val_loss: 17.3459 - val_mse: 19.4792


Epoch 16/50
12/12  0s 9ms/step - loss: 21.9437 - mse: 22.0339 - val_loss: 17.2272 - val_mse: 19.1458


Epoch 17/50
12/12  0s 7ms/step - loss: 21.6555 - mse: 21.7453 - val_loss: 16.9615 - val_mse: 18.7745


Epoch 18/50
12/12  0s 7ms/step - loss: 21.3495 - mse: 21.4411 - val_loss: 16.3717 - val_mse: 18.2899


Epoch 19/50
12/12  0s 7ms/step - loss: 20.6726 - mse: 20.7630 - val_loss: 15.9535 - val_mse: 17.9287


Epoch 20/50
12/12  0s 7ms/step - loss: 20.2702 - mse: 20.3593 - val_loss: 16.0443 - val_mse: 17.9347


Epoch 21/50
12/12  0s 8ms/step - loss: 20.2047 - mse: 20.2942 - val_loss: 16.1980 - val_mse: 18.0644


Epoch 22/50
12/12  0s 8ms/step - loss: 19.9128 - mse: 20.0025 - val_loss: 16.0767 - val_mse: 18.0009


Epoch 23/50
12/12  0s 9ms/step - loss: 19.6004 - mse: 19.6896 - val_loss: 16.2789 - val_mse: 18.1460


Epoch 24/50
12/12  0s 8ms/step - loss: 19.2893 - mse: 19.3797 - val_loss: 16.0885 - val_mse: 17.9629


Epoch 25/50
12/12  0s 9ms/step - loss: 18.7077 - mse: 18.7943 - val_loss: 16.8306 - val_mse: 18.6285


Epoch 26/50
12/12  0s 9ms/step - loss: 18.9461 - mse: 19.0330 - val_loss: 16.4194 - val_mse: 18.2690


Epoch 27/50
12/12  0s 9ms/step - loss: 18.3644 - mse: 18.4507 - val_loss: 16.5490 - val_mse: 18.3886


Epoch 28/50
12/12  0s 9ms/step - loss: 18.2240 - mse: 18.3101 - val_loss: 16.5265 - val_mse: 18.4681


Epoch 29/50
12/12  0s 8ms/step - loss: 17.9240 - mse: 18.0101 - val_loss: 16.4153 - val_mse: 18.4507


Epoch 30/50
12/12  0s 16ms/step - loss: 17.6055 - mse: 17.6908 - val_loss: 16.4671 - val_mse: 18.5082


Epoch 31/50
12/12  0s 9ms/step - loss: 17.4538 - mse: 17.5392 - val_loss: 16.2644 - val_mse: 18.3329


Epoch 32/50
12/12  0s 8ms/step - loss: 17.1733 - mse: 17.2598 - val_loss: 16.0222 - val_mse: 18.1793


Epoch 33/50
12/12  0s 9ms/step - loss: 16.9461 - mse: 17.0327 - val_loss: 15.7303 - val_mse: 17.9443


Epoch 34/50
12/12  0s 13ms/step - loss: 16.6271 - mse: 16.7137 - val_loss: 15.5304 - val_mse: 17.8173


Epoch 35/50
12/12  0s 8ms/step - loss: 16.3307 - mse: 16.4184 - val_loss: 14.6386 - val_mse: 16.8968

Epoch 36/50
12/12  0s 7ms/step - loss: 15.9605 - mse: 16.0490 - val_loss: 14.3735 - val_mse: 16.6316

Epoch 37/50
12/12  0s 7ms/step - loss: 15.7633 - mse: 15.8517 - val_loss: 14.3241 - val_mse: 16.6089

Epoch 38/50
12/12  0s 7ms/step - loss: 15.4102 - mse: 15.4995 - val_loss: 13.9847 - val_mse: 16.2224

Epoch 39/50
12/12  0s 8ms/step - loss: 15.0752 - mse: 15.1640 - val_loss: 13.9605 - val_mse: 16.2534

Epoch 40/50
12/12  0s 7ms/step - loss: 15.0411 - mse: 15.1314 - val_loss: 13.0942 - val_mse: 15.3455

```

Epoch 41/50
12/12 ————— 0s 12ms/step - loss: 14.7562 - mse: 14.8464 - val_loss: 13.2051 - val_mse: 15.5548
Epoch 42/50
12/12 ————— 0s 9ms/step - loss: 14.4934 - mse: 14.5852 - val_loss: 12.5080 - val_mse: 14.7156
Epoch 43/50
12/12 ————— 0s 7ms/step - loss: 14.1859 - mse: 14.2789 - val_loss: 11.8898 - val_mse: 13.9533
Epoch 44/50
12/12 ————— 0s 8ms/step - loss: 13.7463 - mse: 13.8391 - val_loss: 11.6927 - val_mse: 13.6303
Epoch 45/50
12/12 ————— 0s 9ms/step - loss: 13.5959 - mse: 13.6895 - val_loss: 11.2100 - val_mse: 13.0022
Epoch 46/50
12/12 ————— 0s 9ms/step - loss: 13.3295 - mse: 13.4239 - val_loss: 10.6677 - val_mse: 12.0708
Epoch 47/50
12/12 ————— 0s 11ms/step - loss: 12.6864 - mse: 12.7808 - val_loss: 10.3908 - val_mse: 11.6074
Epoch 48/50
12/12 ————— 0s 7ms/step - loss: 12.3872 - mse: 12.4816 - val_loss: 10.1643 - val_mse: 11.2407
Epoch 49/50
12/12 ————— 0s 15ms/step - loss: 12.0128 - mse: 12.1064 - val_loss: 9.9178 - val_mse: 10.8233
Epoch 50/50
12/12 ————— 0s 9ms/step - loss: 11.7226 - mse: 11.8161 - val_loss: 9.7806 - val_mse: 10.5111

```

Model Evaluation

```

In [74]: # Preview the mean value of training and validation data
y_train.mean(), y_val.mean()

```

```

Out[74]: (22.235537, 24.89756)

```

```

In [75]: # Evaluate the model on the test data
print("Evaluation on Test data \n")
loss, mse = model.evaluate(X_test, y_test, batch_size=32)
print(f"\nModel loss on test set: {loss}")
print(f"Model mean squared error on test set: {(mse):.2f}")

```

Evaluation on Test data

```

4/4 ————— 0s 5ms/step - loss: 13.4026 - mse: 14.2381

```


Model loss on test set: 12.777583122253418

Model mean squared error on test set: 14.87

```

In [76]: # Plot the Loss curves
pd.DataFrame(history.history).plot(figsize=(6, 4), xlabel="Epochs", ylabel="Loss",
plt.show()

```



Model Prediction

```
In [77]: # Make predictions
y_pred = model.predict(X_test)

# View the first prediction
y_pred[0]
```

4/4 ————— 0s 25ms/step

```
Out[77]: array([19.273684], dtype=float32)
```